



The two models rank the same predictors differently (Spearman rho = 0.37)

A. OLS: one partial effect for everyone



B. SHAP: a contribution per respondent



Panel A: standardized OLS coefficients, so 0.06 means a one-standard-deviation increase is associated with a 6-percentage-point higher informal share, holding every other predictor fixed; bars are 95% confidence intervals. Panel B: each dot is one held-out respondent, horizontal position is how far that feature moved that person's prediction away from the average, and color is the feature value, red high and blue low. A coefficient and a mean absolute SHAP value are not the same quantity, which is why the rankings diverge. The two features that move furthest between the rankings are `lives_alone`, which is first by SHAP and 24th by coefficient because the living-arrangement indicators are nearly collinear, and `total_care_hours`, whose relationship to the outcome is non-monotone.