


Survey weights change the estimand; household clustering turns out not to matter here



HRS 2022, $n = 2,142$. Panel A compares the sample mean with the population-weighted mean using the HRS respondent weight SWGTR (available for 80% of the sample; Kish design effect 1.83). Panel B plots each standardized coefficient unweighted against weighted; points off the diagonal are estimates that depend on which population is being described. Panel C compares two 5-fold designs on the identical model: random splitting of individuals versus splitting whole households. Only 5% of households contribute more than one respondent, so there is little to leak; the difference is smaller than the error bars, which are the standard deviation across folds, and runs in the opposite direction to leakage. GroupKFold is adopted anyway as the defensible default. Note that Panel A drops all 73 nursing-home residents, who have no survey weight, so the weighted column describes the community-dwelling population only.