


**Figure 10. A neural network buys nothing on 2,142 rows of tabular survey data**



HRS 2022,  $n = 2,142$ , 33 features, scored by 5-fold GroupKFold on household so no household spans the train/test boundary. Panels A and B search over architecture and L2 penalty; the flat, noisy response to added capacity in Panel A and the improvement under heavier penalty in Panel B are both symptoms of a model with far more parameters than the data can identify. Error bars in Panel C are the standard deviation across the five folds and are wide enough to overlap, so the ordering between the top families should be read as 'indistinguishable', not as a ranking.